

# Hybrid ML Scheduler - Project Documentation

Version: 2.0.0 Date: November 27, 2025

## 1. Executive Summary

The **Hybrid ML Scheduler** is an advanced simulation and scheduling system designed to optimize task allocation in heterogeneous computing environments (CPU + GPU). It leverages **Machine Learning (Random Forest)** and **Reinforcement Learning (DQN)** to make intelligent scheduling decisions that balance execution time, energy consumption, and cost.

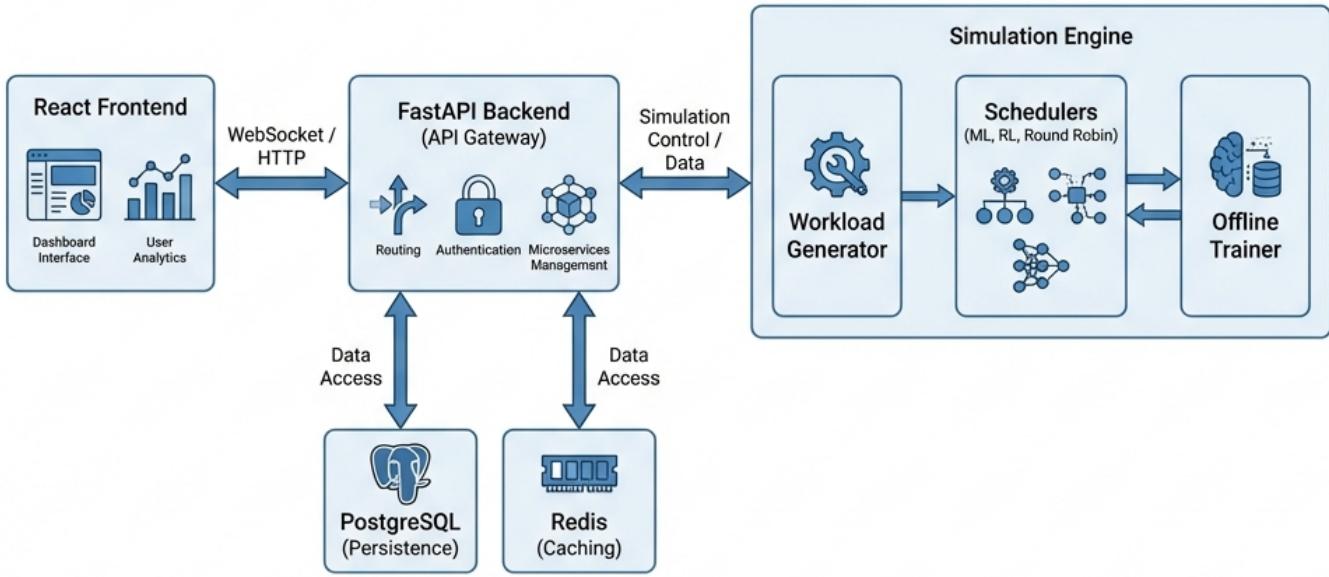
The system features a continuous simulation engine, a robust backend API, persistent storage, and a real-time interactive dashboard for visualization and analysis.

## 2. System Architecture (High-Level Design)

The system follows a modern, modular architecture composed of four main layers:

1. **Presentation Layer (Frontend):** A React-based dashboard for real-time monitoring and control.
2. **Application Layer (Backend):** A FastAPI server handling API requests, WebSocket streaming, and business logic.
3. **Simulation Layer (Engine):** A Python-based engine that generates workloads and executes scheduling strategies.
4. **Data Layer (Persistence):** PostgreSQL for long-term storage and Redis for high-speed caching.

# Hybrid ML Scheduler System Architecture



## Data Flow Overview

- Workload Generation:** The Simulation Engine generates synthetic tasks with varying characteristics (size, compute intensity, memory).
- Scheduling:** Tasks are processed by multiple schedulers in parallel (Round Robin, Random, Greedy, Hybrid ML, RL Agent, Oracle).
- Execution Simulation:** A Virtual Cluster model estimates execution time and energy based on the scheduling decision.
- Data Persistence:** Results are saved to PostgreSQL via the Backend API. Training data is buffered and stored.
- Model Retraining:** The Offline Trainer periodically retrains the ML model using the latest historical data from the database.
- Visualization:** The Backend broadcasts real-time updates via WebSockets to the Frontend Dashboard.

# 3. Deep Dive: Core Logic & Algorithms

This section explains the internal mechanics of the system, allowing you to understand the "how" and "why" without reading the code.

## 3.1. Workload Generation (The "Tasks")

The system generates a continuous stream of synthetic tasks that mimic real-world parallel computing jobs.

- **Generation Process:** Tasks arrive according to a **Poisson Process** (exponentially distributed inter-arrival times), creating a realistic, bursty workload.

### Task Attributes:

- **Size (\$N\$):** The magnitude of the problem (100 - 5000 units).

**Compute Intensity (\$I\$):** A value between 0.0 and 1.0 indicating how much the task benefits from parallelization.

- $I \approx 1.0$ : Highly parallelizable (Matrix Multiplication, Deep Learning).
- $I \approx 0.0$ : Serial (I/O bound, recursive logic).

- **Memory Required (\$M\$):** RAM usage (10 - 500 MB).

- **Duration Model:** The estimated base duration is calculated as:  $T_{\text{base}} \propto \frac{N^{1.5}}{I + 0.5}$  *This means larger tasks take super-linearly longer, but high compute intensity reduces time (assuming parallel hardware).*

## 3.2. The Schedulers (The "Competitors")

The system runs six scheduling strategies in parallel for every task to compare their performance.

### 1. Round Robin (Baseline)

- **Logic:** Alternates blindly between resources.
- **Behavior:** Task  $i$  goes to GPU, Task  $i+1$  goes to CPU.
- **Pros/Cons:** Simple but inefficient; sends GPU-hostile tasks to GPU and vice versa.

### 2. Random (Baseline)

- **Logic:** Assigns a random fraction of the task to the GPU (0.0 to 1.0).
- **Pros/Cons:** Acts as a stochastic baseline to prove that other methods are learning.

### 3. Greedy (Heuristic)

- **Logic:** Uses the task's **Compute Intensity** directly as the GPU fraction.
- **Formula:**  $\text{Fraction}_{\text{GPU}} = \text{Intensity}$
- **Rationale:** High intensity tasks *should* go to GPU. This is a strong heuristic baseline.

### 4. Hybrid ML (The "Brain")

- **Type:** Supervised Learning (Random Forest Regressor).

#### Features:

- Raw: Size, Intensity, Memory.
- Derived: Memory per Unit Size, Compute to Memory Ratio.

#### Logic:

1. Predicts the **Optimal GPU Fraction** ( $y$ ) that minimizes execution time.

- Selects the specific GPU ID that minimizes a weighted cost of **Time** and **Energy**.
- Training:** Retrains every 50 tasks using a sliding window of the last 1000 "Oracle" decisions.

## 5. RL Agent (Deep Q-Network)

- Type:** Reinforcement Learning (DQN with Dueling Architecture).
- State:** Vector  $[\text{Size}, \text{Intensity}, \text{Memory}]$ .
- Action Space:** Discrete choices:  $\{\text{CPU}, \text{GPU}_0, \text{GPU}_1, \text{GPU}_2, \text{GPU}_3\}$ .
- Reward Function:** Negative weighted sum of Time and Energy.  $R = - (0.5 \times \text{Time} + 0.5 \times \text{Energy})$
- Learning:** Uses "Experience Replay" to learn from past mistakes and optimizes its policy to maximize long-term rewards.

## 6. Oracle (The "Ground Truth")

- Logic:** A theoretical solver that "cheats" by trying every possible split (0% to 100% in 5% steps).
- Purpose:** It finds the absolute mathematical minimum execution time for a task.
- Usage:**
  - Acts as the **Label** for the Hybrid ML model (Supervised Learning).
  - Serves as the **Performance Ceiling** (100% Efficiency) for comparison.

## 3.3. Simulation Physics

How do we calculate "Time" and "Energy"?

### Execution Time:

- CPU Time:** Base duration.
- GPU Time:**  $\frac{\text{Base Duration}}{\text{Speedup}} + \text{TransferTime}$
- Speedup:**  $1.0 + (3.0 \times \text{Intensity})$  (Max 4x speedup for high intensity).
- Transfer Time:**  $\frac{\text{Memory}}{\text{Bandwidth}}$  (Simulating PCIe bottlenecks).

### Energy Consumption:

- GPU Power:** 50W (Active).
- CPU Power:** 30W (Active).
- $\text{Energy} (\text{Joules}) = \text{Power} \times \text{Time}$ .

## 4. System Flow & Architecture

### 4.1. The "Loop"

- Generate:** A new task is born (`WorkloadGenerator`).
- Broadcast:** The task is sent to all 6 schedulers simultaneously.
- Decide:** Each scheduler makes its move (Predict, Randomize, or Calculate).
- Simulate:** The `VirtualMultiGPU` calculates the *result* (Time, Energy) for each decision.
- Persist:** Results are saved to PostgreSQL.
- Learn:**
  - Hybrid ML:** If 50 tasks have passed, fetch history -> Retrain Random Forest.
  - RL Agent:** Store transition -> Update Q-Network weights.
- Visualize:** Send JSON packet via WebSocket to the Dashboard.

## 4.2. Component Details

### Frontend Dashboard (/dashboard)

- **Tech:** React, Vite, TailwindCSS, Recharts.

#### Key Views:

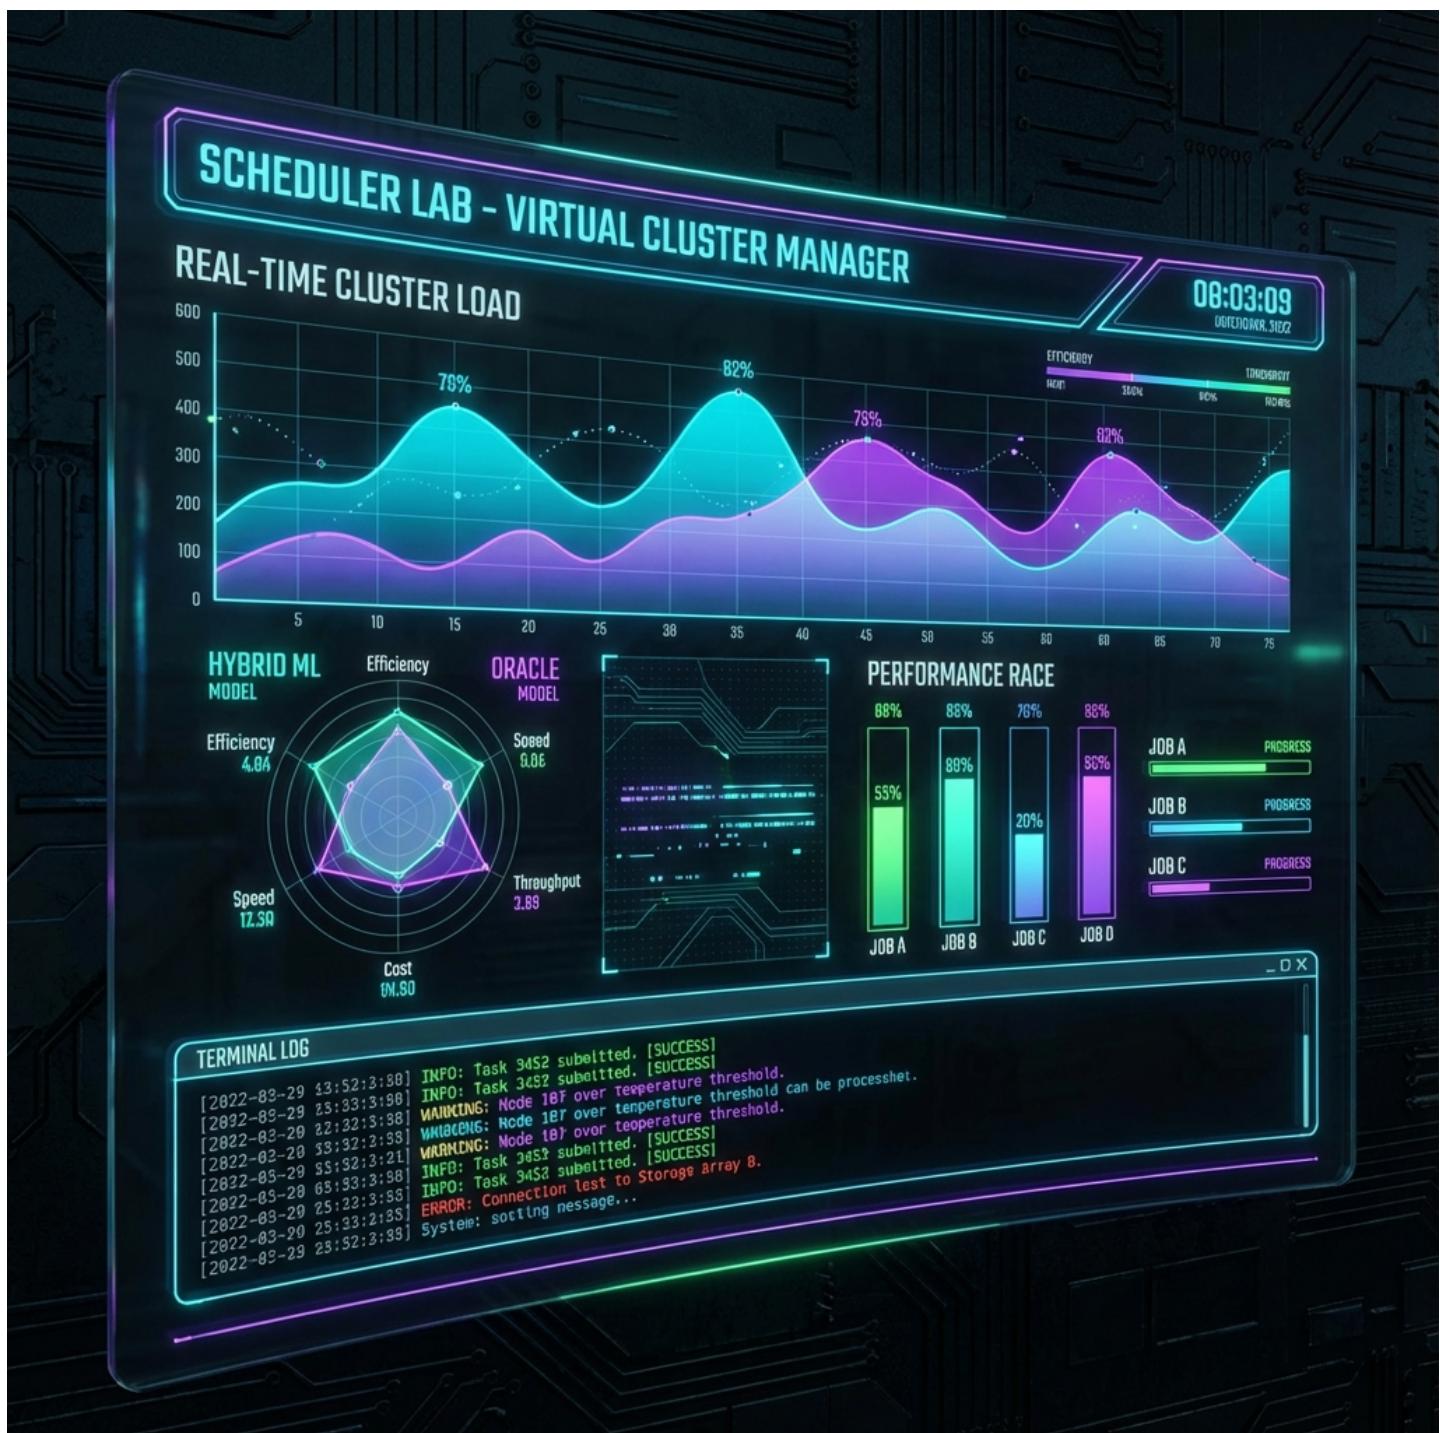
- **Performance Race:** A live bar chart where shorter bars = faster schedulers.

#### Enhanced Analytics:

- **Heatmap:** Shows which tasks (Size vs. Intensity) perform best on GPU.

- **Win/Loss Matrix:** How often Scheduler A beats Scheduler B.

- **State Management:** Uses React `useState` and `useEffect` for real-time data updates.



### Backend Server (/backend)

- **Tech:** FastAPI, Uvicorn, SQLAlchemy (Async), Pydantic.
- **Role:** The central nervous system. It orchestrates the simulation, manages the DB connection, and serves the API.
- **Security:** Implements Rate Limiting (Redis) and Input Validation to protect the system.

## Data Layer

- **PostgreSQL:** Stores the "Truth". Every single task execution is logged here.
- **Redis:** The "Short-term Memory". Caches high-speed data like current stats to prevent DB overload.

# 5. Technical Reference

## 5.1. Database Schema (PostgreSQL)

The system uses a relational schema optimized for time-series performance.

### `tasks` Table

Stores metadata for every generated task. | Column | Type | Description | |-----|-----|-----| | task\_id | Integer (PK) | Unique identifier for the task. || size | Float | Problem size (\$N\$). || compute\_intensity | Float | Parallelizability factor (\$0.0 - 1.0\$). || memory\_required | Float | RAM required in MB. || arrival\_time | Float | Simulation timestamp of arrival. || dependencies | JSONB | List of parent task IDs. |

### `scheduler_results` Table

Logs the outcome of every scheduling decision. | Column | Type | Description | |-----|-----|-----| | id | Integer (PK) | Unique result ID. || task\_id | Integer (FK) | Reference to the task. || scheduler\_name | String | Name of the strategy (e.g., 'hybrid\_ml'). || gpu\_fraction | Float | Allocated GPU portion (\$0.0 - 1.0\$). || actual\_time | Float | Execution time in seconds. || energy\_consumption | Float | Energy used in Joules. || execution\_cost | Float | Cost in USD. |

### `training_data` Table

Historical data used to train the Hybrid ML model. | Column | Type | Description | |-----|-----|-----| | size | Float | Task size. || compute\_intensity | Float | Task intensity. || optimal\_gpu\_fraction | Float | **Label:** The best fraction found by Oracle. || optimal\_time | Float | The execution time achieved by Oracle. |

## 5.2. API Specification (FastAPI)

### Simulation Control

- `POST /api/simulation/start`: Begin the continuous simulation.
- `POST /api/simulation/stop`: Gracefully stop the engine.
- `POST /api/simulation/pause`: Temporarily halt task generation.
- `GET /api/simulation/status`: Get current stats (tasks processed, running state).

### Data & Metrics

- `GET /api/full_history`: Retrieve the latest 1000 training records (used for charts).
- `GET /api/metrics`: Expose Prometheus-formatted metrics for scraping.
- `GET /api/health`: Check connectivity to PostgreSQL and Redis.

## WebSocket (/ws)

- **Protocol:** JSON over WebSocket.

### Events:

- `simulation_update`: Real-time packet with current task, scheduler results, and cluster utilization.
- `notification`: System alerts (e.g., "Model Retrained").

## 5.3. Configuration Management

The system is configured via environment variables (using `pydantic-settings`).

### Key Variables (.env)

Variable	Default	Description
ENVIRONMENT	development	App environment (dev/prod).
POSTGRES_HOST	localhost	Database host.
POSTGRES_DB	hybrid_scheduler_db	Database name.
REDIS_HOST	localhost	Redis cache host.
NUM_GPUS	4	Number of virtual GPUs to simulate.
RETRAIN_INTERVAL	50	Tasks between model updates.

## 6. Directory Structure

```
hybrid_ml_scheduler/ ├── backend/ # FastAPI Application └── api/ # Routes and
                      # Controllers └── core/ # Config and DB setup └── middleware/ # Rate Limit,
                      # Security └── models/ # SQLAlchemy & Pydantic models └── services/ # Business
                      # Logic (Data, Cache) └── dashboard/ # React Frontend └── src/ └── components/
                      # Reusable UI widgets (Charts) └── App.jsx # Main Dashboard View └── src/ # Simulation
                      # Engine └── simulation_engine.py # Main Loop └── online_scheduler.py # Hybrid ML Logic
                      └── dqn_scheduler.py # RL Logic └── workload_generator.py # Task Factory
                      └── scripts/ # Utility Scripts (Init DB, Verify) └── tests/ # Unit and
                        # Integration Tests
```